

Dear Code Metal hiring team,

Code Metal's pitch is one I've already tested at a smaller scale. When I rolled out Claude Code across my team at Pharmacists Mutual, I didn't trust the model to grade its own homework. I built a security framework instead: least-privilege SQL logins, managed settings, a custom MCP server - the works. That framework cut our Agile delivery time roughly in half.

I think I fit the AI Agentic Product path you're hiring for. I've pushed agentic coding into my org for both the data and engineering teams. On the Python side, I've built and maintain a running FastAPI backend for dyingskies.com, and use the spaCy library when ingesting raw text files at gufime.com. My Dagster asset graphs for Gutenberg-Fingerprint run a staged pipeline like your post describes: parse, transform, validate, with a sensor that skips the run when there's nothing new.

Right now I lead a small data engineering team. I own architecture and hiring, and I rebuilt our reporting platform from scratch, swapping opaque vendor logic for documented, tested systems. That cut storage 50% and processing time 60%. I've seen what it costs an organization to run on code nobody trusts, and I want my next role to involve fixing that at the source.

Your website and post mentioned a few "modest" customers: the Air Force, L3Harris, Toshiba. High stakes - it's obvious why your code has to be provably correct. I'd be pleased as punch for a chance to chat about how I could contribute to Code Metal's product.

All the best,
Sander VanWilligen