

I'm going to answer this question, but I'm going to try a different tactic on this app. I've never tried begging, so I'm begging you: please, help me escape from this ten-year, impossibly-slow insurance hell! I can't stand it, I'm losing my mind. Every decision needs a triple confirmation, everyone sits in silence in the meetings, until one person speaks up and repeats what the last person said. We "can't adopt that fast, quick-turnaround product, because we have our legacy monolith system we need to support," or "We can't allow you to use Claude-Code SQL-server MCPs, because we're concerned about Claude-Code reading our data." Please, help - I can't take it anymore!

All of my work is build with AI agents, specifically Claude-Code and openRouter/opencode. One of my favorites was the split-duties framework I build for dying-skies (<https://github.com/sandervw/dying-skies>). I basically had a monorepo with four parts, where I had: 1 developer agent per part, 1 code-standards/security-validator per part, and one integration tester between parts. For the code-standards one I hand-tuned the agent (no comments above 12 characters, no files longer than 300 lines, stuff like that). The ending-code didn't just work, but was also super-clean and understandable.

Please - help a diligent, AI-integrated engineer out!

- Sander VanWilligen