

Dear Tailscale hiring team,

I'm applying for the Analytics Engineer, Data role. Your post asks for someone who brings clarity to chaos and builds away manual work. That's been my last decade. In 2020 the analytical reporting vendor at my company went under, and I helped start the replacement. Then I took over the ETL, the SQL architecture, and the Tabular/Power BI stack. In 2024 I tore the enterprise reporting down and rebuilt it from scratch, swapping opaque vendor logic for documented SSIS and SQL. Storage dropped 50%. Processing time dropped over 60%.

For years I was the sole owner of data architecture and strategy, reporting straight to company leadership. The numbers had to hold up. I rerouted ingestion through Ataccama MDM instead of pulling from source, and I ran the evaluation that picked it. Since March 2026 I'm Lead Data Engineer, which means hiring and onboarding. I trained the team on our ETL and MDM platform. Before that, one person understood it. Now everyone can run it, and writing things down is bread and butter for me.

dbt is where all my latest personal work lives. My day job is SQL Server and SSIS, mid-migration to a medallion lakehouse on Azure and Microsoft Fabric. So I built Gutenberg-Fingerprint on my own, a nightly change-data-capture pipeline I run solo on a ~\$5/month VPS. Postgres, dbt Core, and Dagster, plus a fact constellation, SCD2 snapshots, source freshness checks, audit logging, and a full test suite. It compiles against DuckDB in dev and Postgres in prod. Ingestion to visualization, the works.

I'd be happy to build the models that your Finance/Engineering/Sales departments digest. Quick fixes and long-term design both sound fine to me. I haven't touched Snowflake, Orb, or Stripe yet. The modeling and validation habits come with me, so I expect to catch up fast. I'm ready to leave insurance for a company whose product I'd use myself.

All the best,
Sander VanWilligen