

Dear Trulioo hiring team,

I'm applying for the Staff Engineer or Architect role on the Agentic AI team. Your post calls for agents that hold up in regulated work, with evidence behind every finding. That's the work I want to be doing.

I've shipped prod systems in TypeScript, Python, and SQL for ten years. As Lead Data Engineer at Pharmacists Mutual, I ran the evaluation of agentic coding tools, then wrote the security framework that cleared them for prod. Guardrails, approval flows, a custom SQL MCP server, the works. Auditability and human sign-off went in from day one. I use agents in my own daily engineering too, and I still own whatever they hand back.

I like defining the problem before I solve it. As sole technical decision-maker, reporting to company leadership, I replaced opaque vendor logic with a documented, tested reporting platform. Storage dropped 50% and processing time dropped over 60%. My first job was writing a nightly Java process for government OFAC compliance screenings, so compliance-heavy screening is familiar ground.

I work with the Claude and OpenAI APIs and build MCP servers. Agent frameworks and evaluation harnesses are where I'm thinner, and I'd rather say so up front - but I'm not unfamiliar, and I learn fast. I'm comfortable working remotely within North America.

Happy to talk if it sounds like I might be a fit for the role.

All the best,
Sander VanWilligen