

Dear G2 hiring team,

Your post asks for engineers who use agent harnesses daily. I've been doing that for a little over two years now. I championed and piloted Claude Code across the data and software engineering teams at my current company. To do that, I wrote the security framework that got it cleared for production use (least-privilege SQL logins, approval flows, a custom MCP server) - the results of that pilot cut our Agile story delivery time roughly in half. (Somewhat selfishly) I also just enjoy building with AI and Large Language Models - I think we're in a golden age of tech.

Full-stack is my bread and butter. I've spent the last ten years building from tip-to-tail: React/Astro front ends, Node/Express and FastAPI services, hand-rolled auth and rate limiting, Postgres, Mongo, SQL Server, Terraform and CI, and Cloudflare/Azure hosting. My dyingskies.com is just one example - four services in a monorepo, frontend, backend, infra, and Observable for the Analytics. gufime.com is another, a nightly Dagster pipeline, loading uncleaned, raw text data, parsing and analyzing it using Python's spaCy library. I run it solo on a \$5/month VPS (I ran it on Azure initially, but I'm cost-conscious - a VPS does the job for 1/2 the code and cost), with sensors and tests keeping it honest.

I've built directly with frontier models too. Claude API and SDK work, OpenAI API, MCP servers, agentic workflows baked into how I write code every day. I've been using the LLM APIs and agents since they dropped - I stay up-to-date by reading the Anthropic/OpenRouter/Hermes-agent blogs.

So, if you're looking for someone who (a) loves using agents and AI, and (b) knows the full stack techs and how to leverage AI to write them, I'd welcome the chance to talk about the role.

All the best,
Sander VanWilligen