

Dear Render Data team,

I want to help Render conquer the challenges described in your posting: platform discovery, turning deployment into adoption, scaling, finding the actual dollar-makers. I've spent the last 9 years of my data-driven career solving those exact same problems (in the much-less exciting field of Insurance); I'd be pleased as punch to solve them for an AI-native cloud company.

I currently own analytics end to end for a specialty insurer, reporting directly to upper management as the sole data strategist and Lead Data Engineer. Alongside the actual building of pipelines, that meant deciding what to measure, defining metrics leadership actually uses to generate dollars, turning ambiguous executive questions into fact-based analysis, and - ultimately - rebuilding our entire analytical platform from scratch. End result? A documented model that cut storage 50%, processing time over 60 percent.

The analytics-engineering foundation you describe is my precise work today, professionally and personally. I build dbt Core models with SCD2 snapshots, source freshness, full test suites, Dagster orchestration, dual-target compilation against DuckDB and Postgres, and Fabric, Evidence.dev, and observable semantic layers. My side projects, Gutenberg-Fingerprint and DyingSkies, run that modern-data stack solo on a five-dollar/month VPS, deployed to Cloudflare.

I am also the person who brought agentic-assisted development into my org. I ran the Claude Code evaluation, authored the security framework, and cut agile story delivery time by another 60%. Your call for people who use Claude Code to accelerate analysis is my daily bread and my abiding butter.

All the best,
Sander VanWilligen