

Dear Render hiring team,

Render's Data Platform role is my kind of problem - set the architecture, then go live in prod. I lead the data engineering team at Pharmacists Mutual, where I'm setting the direction for a migration off on-prem SQL Server and SSIS onto a medallion lakehouse on Azure and Microsoft Fabric. Your post's mix of warehouse ownership, orchestration, architectural direction, and mentoring lines up with my existing daily-do's.

At Pharmacists Mutual I own warehouse reliability end to end - query performance, cost, and access control, the works - through Ataccama MDM/DQ and final analytical reporting to senior leadership. I rebuilt our ETL and Tabular/Power BI reporting from scratch, cut storage 50%, and cut processing time over 60%. I also recruited and onboarded a senior data engineer, and turned pipeline knowledge that used to live in one person's head into something the whole team owns.

Outside work I run Gutenberg-Fingerprint solo on a \$5/month VPS - Terraform/Tofu, Postgres, dbt Core, and Dagster OSS orchestrating a CDC pipeline with SCD2 snapshots and source freshness checks. I handle the CI/CD with GitHub Actions.

I also led Claude Code adoption at Pharmacists Mutual. I wrote the security framework that got it approved for prod, and story delivery time dropped roughly in half. Somewhat selfishly, I admit, I like tools that make the boring parts of the job faster.

I'd welcome a conversation about the role.

All the best,
Sander VanWilligen