

Dear Verantos hiring team,

I'm a Lead Data Engineer with ten years of building and owning prod pipelines for messy, high-stakes data. I want to bring that to a team putting out real-world clinical data for regulators and biopharma, especially as you keep expanding sources and integrations the way you described in your posting.

At Pharmacists Mutual, I lead the team responsible for our analytical data platform. We're currently mid-migration from SQL Server/SSIS to a medallion lakehouse on Azure and Fabric. I redesigned our ingestion pipeline to route through Ataccama MDM so data quality issues get caught before they reach downstream reporting. I'm also pushing for investment and learning of dbt, Dagster, and Python at my company - somewhat selfishly, I admit. On my own time I built Gutenberg-Fingerprint, a CDC pipeline with raw text Python ingestion (using spaCy and Pandas), a dbt fact constellation, SCD2 snapshots, source freshness checks, a full test suite, a nightly Dagster orchestration, running unattended on a VPS that costs me ~\$5/month.

I also led our Claude Code rollout, after crusading for agentic adoption for our engineering department. Previously, business thought 'agentic' meant Microsoft Chat in Teams. I wrote the security framework that got agentic AI tooling approved for prod use, which ended up cutting our Agile delivery time roughly in half. I use AI daily in my own work, and like bringing other engineers up to speed as well.

I don't have prod Snowflake experience, but my dbt and SQL background transfers, and I pick up platform specifics fast. I'd welcome the chance to talk about the role.

All the best,
Sander VanWilligen